

Assignment 1

ELEC 576 - Prof. Patel

Adrian E. Radillo – S01254516

October 5, 2017

Contents

1	Backpropagation in a simple neural network	1
1.a	Dataset	1
1.b	Activation function	1
1.c	Build the neural network	2
1.d	Backward pass - backpropagation	3
1.e	Time to have fun - training	5
1.f	Even more fun - training a deeper network!!!	6
2	Training a simple Deep Convolutional Network on MNIST	11
2.a	Build and train a 4-layer DCN	11
2.a.1	Build network	11
2.a.2	Set up training	12
2.a.3	Run training	14
2.a.4	Visualize training	15
2.b	More on visualizing your training	15
2.c	Time for more fun!	15

GitHub: The Python code used for this project is available at the public repository, <https://github.com/aernesto/IntroMachineLearningCourse.git>

1 Backpropagation in a simple neural network

1.a Dataset

1.b Activation function

Our implementation of actFun is:

```
1 def actFun(self, z, type):
2     '''
3     actFun computes the activation functions
4     :param z: net input
5     :param type: Tanh, Sigmoid, or ReLU
6     :return: activations
7     '''
8     if type == 'Tanh':
9         return np.tanh(z)
10    elif type == 'Sigmoid':
11        return 1. / (1 + np.exp(-z))
12    elif type == 'ReLU':
13        return np.maximum(z, np.zeros(z.shape))
```

The derivatives of the three activation functions (denoted by f) are:

Tanh $f'(z) = 1 - (f(z))^2$

Sigmoid $f'(z) = f(z)(1 - f(z))$

```
if __name__ == "__main__":
    main()
```




Figure 1: Data from Make-Moons dataset

ReLU

$$f'(z) = \begin{cases} 0 & \text{if } z < 0 \\ \text{undefined} & \text{if } z = 0 \\ 1 & \text{if } z > 0 \end{cases}$$

Our implementation of `diff_actFun` therefore is:

```
1 def diff_actFun(self, z, type):
2     '''
3     diff_actFun computes the derivatives of the activation functions wrt the net
4     input
5     :param z: net input
6     :param type: Tanh, Sigmoid, or ReLU
7     :return: the derivatives of the activation functions wrt the net input
8     '''
9     if type == 'Tanh':
10         return 1 - (self.actFun(z, type))**2
11     elif type == 'Sigmoid':
12         return self.actFun(z, type) * (1 - self.actFun(z, type))
13     elif type == 'ReLU':
14         return (z > 0).astype(float)
```

1.c Build the neural network

Our implementation of `feedforward` and `calculate_loss` is:

```
1 def feedforward(self, X, actFun_2):
2     '''
3     feedforward builds a 3-layer neural network and computes the two
4     probabilities,
5     one for class 0 and one for class 1
6     :param X: input data
7     :param actFun: activation function
8     :return:
9     '''
10     self.z1 = X.dot(self.W1) + np.tile(self.b1, (X.shape[0],1)) # dim 200 x nH
11     self.a1 = actFun_2(self.z1)
12     self.z2 = self.a1.dot(self.W2) + np.tile(self.b2, (self.a1.shape[0],1))
13     exp_scores = np.exp(self.z2)
```

```

13         self.probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
14         return None

1 def calculate_loss(self, X, y):
2     """
3     calculate_loss computes the loss for prediction
4     :param X: input data
5     :param y: given labels
6     :return: the loss for prediction
7     """
8     num_examples = len(X)
9     self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
10
11     # Calculating the loss
12     y_hot = np.array([[1, 0] if yy == 0 else [0, 1] for yy in y])
13     data_loss = -num_examples * np.sum(y_hot * np.log(self.probs))
14
15     # Add regulatization term to loss (optional)
16     data_loss += self.reg_lambda / 2 * (np.sum(np.square(self.W1)) + np.sum(np.
17         square(self.W2)))
18     return (1. / num_examples) * data_loss

```

1.d Backward pass - backpropagation

This is our implementation of backprop:

```

1 def backprop(self, X, y):
2     """
3     backprop implements backpropagation to compute the gradients used to update
4     the parameters in the backward step
5     :param X: input data
6     :param y: given labels
7     :return: dL/dW1, dL/b1, dL/dW2, dL/db2
8     """
9     num_examples = len(X)
10    delta3 = self.probs
11    delta3[range(num_examples), y] -= 1
12    dW2 = np.zeros(self.W2.shape)
13    db2 = np.zeros(self.b2.shape)
14    dW1 = np.zeros(self.W1.shape)
15    db1 = np.zeros(self.b1.shape)
16
17    for example in range(num_examples):
18        dW2 += np.multiply(np.tile(delta3[example, :], (self.nn_hidden_dim, 1)),
19            np.tile(self.a1[example, :], (self.nn_output_dim, 1)).T)
20        db2 += delta3[example, :]
21        dW1 += np.tile(np.sum(np.tile(delta3[example, :],
22            (self.nn_hidden_dim, 1)) * self.W2,
23            axis = 1), (self.nn_input_dim, 1)) * \
24            np.tile(self.diff_actFun(self.z1[example, :], self.actFun_type),
25            (self.nn_input_dim, 1)) * \
26            np.tile(X[example, :], (self.nn_hidden_dim, 1)).T
27        db1 += np.sum(np.tile(delta3[example, :], (self.nn_hidden_dim, 1)) * self.
28            W2, axis = 1) * \
29            self.diff_actFun(self.z1[example, :], self.actFun_type)
30
31    dW2 /= num_examples
32    db2 /= num_examples
33    dW1 /= num_examples
34    db1 /= num_examples

```

33

34

```
return dW1, dW2, db1, db2
```

1.e Time to have fun - training

We present below the results of training our network with three distinct activation functions, always with 3 hidden nodes:

```
Loss after iteration 0: 0.582912
Loss after iteration 1000: 0.332595
Loss after iteration 2000: 0.311568
Loss after iteration 3000: 0.306815
Loss after iteration 4000: 0.304534
Loss after iteration 5000: 0.303069
Loss after iteration 6000: 0.301871
Loss after iteration 7000: 0.300608
Loss after iteration 8000: 0.299103
Loss after iteration 9000: 0.297249
Loss after iteration 10000: 0.294981
Loss after iteration 11000: 0.292257
Loss after iteration 12000: 0.289076
Loss after iteration 13000: 0.285479
Loss after iteration 14000: 0.281548
Loss after iteration 15000: 0.277399
Loss after iteration 16000: 0.273161
Loss after iteration 17000: 0.268958
Loss after iteration 18000: 0.264898
Loss after iteration 19000: 0.261058
```

```
Loss after iteration 0: 0.622949
Loss after iteration 1000: 0.327359
Loss after iteration 2000: 0.304278
Loss after iteration 3000: 0.300512
Loss after iteration 4000: 0.298879
Loss after iteration 5000: 0.298062
Loss after iteration 6000: 0.297413
Loss after iteration 7000: 0.296828
Loss after iteration 8000: 0.296323
Loss after iteration 9000: 0.295605
Loss after iteration 10000: 0.294933
Loss after iteration 11000: 0.294127
Loss after iteration 12000: 0.292910
Loss after iteration 13000: 0.291803
Loss after iteration 14000: 0.290787
Loss after iteration 15000: 0.289947
Loss after iteration 16000: 0.289227
Loss after iteration 17000: 0.288511
Loss after iteration 18000: 0.287823
Loss after iteration 19000: 0.287297
```

```
Loss after iteration 0: 0.656440
Loss after iteration 1000: 0.525975
Loss after iteration 2000: 0.430989
Loss after iteration 3000: 0.382175
Loss after iteration 4000: 0.359215
Loss after iteration 5000: 0.347746
Loss after iteration 6000: 0.341537
Loss after iteration 7000: 0.337960
Loss after iteration 8000: 0.335806
Loss after iteration 9000: 0.334462
Loss after iteration 10000: 0.333597
Loss after iteration 11000: 0.333025
Loss after iteration 12000: 0.332632
Loss after iteration 13000: 0.332354
Loss after iteration 14000: 0.332148
Loss after iteration 15000: 0.331990
Loss after iteration 16000: 0.331864
Loss after iteration 17000: 0.331761
Loss after iteration 18000: 0.331673
Loss after iteration 19000: 0.331598
```




Figure 2: ‘Tanh’ activation function and 3 hidden nodes




Figure 3: ‘ReLU’ activation function and 3 hidden nodes




Figure 4: ‘Sigmoid’ activation function and 3 hidden nodes

The ‘Tanh’ activation function seems to result in the most flexible decision boundary. It also produces the smallest loss at the end of the training. The ‘Sigmoid’ activation function seems to be the worst of the three, in terms of decision boundary, which is a simple line, and final loss.

When retraining the network with ‘Tanh’ and 50 hidden nodes, we obtain the following output:

```
Loss after iteration 0: 0.506189
Loss after iteration 1000: 0.297850
Loss after iteration 2000: 0.284993
Loss after iteration 3000: 0.276774
Loss after iteration 4000: 0.270326
Loss after iteration 5000: 0.264630
Loss after iteration 6000: 0.258983
Loss after iteration 7000: 0.252944
Loss after iteration 8000: 0.246327
Loss after iteration 9000: 0.239181
Loss after iteration 10000: 0.231729
Loss after iteration 11000: 0.224278
Loss after iteration 12000: 0.217131
Loss after iteration 13000: 0.210523
Loss after iteration 14000: 0.204597
Loss after iteration 15000: 0.199411
Loss after iteration 16000: 0.194953
Loss after iteration 17000: 0.191166
Loss after iteration 18000: 0.187975
Loss after iteration 19000: 0.185300
```




Figure 5: ‘Tanh’ activation function and 50 hidden nodes

In comparing figures 2 and 5, we notice that augmenting the dimension of the hidden layer has two main effects. Firstly, it reduces the final loss, and secondly, it produces a decision boundary which has more degrees of freedom.

1.f Even more fun - training a deeper network!!!

Our full new code is presented here:

```

1 class NeuralNetwork(object):
2     """
3     This class builds and trains a neural network
4     """
5     def __init__(self, layer_sizes, num_examples, actFun_type='Tanh', reg_lambda
        =0.01, seed=0):
6         '''
7         :param layer_sizes: list of positive integers.
8         :param actFun_type: type of activation function. 3 options: 'tanh', 'sigmoid
9         ', 'relu'
10        :param reg_lambda: regularization coefficient
11        :param seed: random seed
12        '''
13
14        '''
15        !!NOTE!! In this code, layer numbering is inverted (this is hidden to the
16        user when he
17        instantiates the class though), therefore the first integer of the list
18        attribute self.layer_sizes
19        is dim of output layer and last integer is dimension of the input layer.
20        '''
21        self.layer_sizes = list(reversed(layer_sizes))
22        self.num_layers = len(self.layer_sizes)
23
24        # Instantiate all the layers, except for the input which contains no
25        # intrinsic weights.
26        layers = {}
27        for n in np.arange(1, self.num_layers):
28            layers[str(n)] = Layer(layer_size = self.layer_sizes[n - 1],
29                                  prec_layer_size = self.layer_sizes[n],
30                                  num_examples = num_examples,
31                                  seed = seed, depth = n)
32
33        # Instantiate input layer separately since no preceding layer
34        layers[str(self.num_layers)] = Layer(layer_size = self.layer_sizes[-1],
35                                              prec_layer_size = 0,
36                                              num_examples = num_examples, seed =
37                                              seed,
38                                              depth = self.num_layers)
39
40        self.layers = layers
41        self.actFun_type = actFun_type
42        self.reg_lambda = reg_lambda
43
44        def actFun(self, z, type):
45            '''
46            actFun computes the activation functions
47            :param z: net input
48            :param type: Tanh, Sigmoid, or ReLU
49            :return: activations
50            '''
51            if type == 'Tanh':
52                return np.tanh(z)

```

```

46         elif type == 'Sigmoid':
47             return 1. / (1 + np.exp(-z))
48         elif type == 'ReLU':
49             return np.maximum(z, np.zeros(z.shape))
50
51     def diff_actFun(self, z, type):
52         """
53         diff_actFun computes the derivatives of the activation functions wrt the net
54         input
55         :param z: net input
56         :param type: Tanh, Sigmoid, or ReLU
57         :return: the derivatives of the activation functions wrt the net input
58         """
59         if type == 'Tanh':
60             return 1 - (self.actFun(z, type))**2
61         elif type == 'Sigmoid':
62             return self.actFun(z, type) * (1 - self.actFun(z, type))
63         elif type == 'ReLU':
64             return (z > 0).astype(float)
65
66     def feedforward(self, X, actFun_2):
67         """
68         feedforward builds a 3-layer neural network and computes the two
69         probabilities,
70         one for class 0 and one for class 1
71         :param X: input data
72         :param actFun: activation function
73         :return:
74         """
75         self.layers[str(self.num_layers)].activations = X
76         for n in reversed(np.arange(1, self.num_layers)):
77             self.layers[str(n)].feedforward(inputs = self.layers[str(n + 1)].
78                                             activations,
79                                             actFun = actFun_2)
80
81         # correct for activations of output layer
82         exp_scores = np.exp(self.layers['1'].linear_output)
83         self.probs = exp_scores / np.sum(exp_scores, axis=1, keepdims=True)
84         self.layers['1'].activations = self.probs
85         return None
86
87     def calculate_loss(self, X, y):
88         """
89         calculate_loss computes the loss for prediction
90         :param X: input data
91         :param y: given labels
92         :return: the loss for prediction
93         """
94         num_examples = len(X)
95         self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
96         # self.feedforward(X, self.actFun) # WHY NOT THIS?
97         # Calculating the loss
98         y_hot = np.array([[1, 0] if yy == 0 else [0, 1] for yy in y])
99         data_loss = -np.sum(y_hot * np.log(self.probs))
100
101         # Add regulatization term to loss (optional)
102         coef = 0
103         for L in np.arange(1, self.num_layers):
104             coef += np.sum(np.square(self.layers[str(L)].weights))

```

```

102     data_loss += self.reg_lambda / ((self.num_layers - 1) * coef)
103     return (1. / num_examples) * data_loss
104
105 def predict(self, X):
106     '''
107     predict infers the label of a given data point X
108     :param X: input data
109     :return: label inferred
110     '''
111     self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
112     return np.argmax(self.probs, axis=1)
113
114 def backprop(self, X, y):
115     '''
116     backprop implements backpropagation to compute the gradients used to update
117     the parameters in the backward step
118     :param X: input data
119     :param y: given labels
120     :return: dL/dW1, dL/b1, dL/dW2, dL/db2
121     '''
122     num_examples = len(X)
123     delta3 = self.probs
124     delta3[range(num_examples), y] -= 1
125     for L in np.arange(1, self.num_layers):
126         if L == 1:
127             self.layers[str(L)].backprop(lambda x: self.diff_actFun(x, type=self
128                                     .actFun_type),
129                                     prev_activations = self.layers[str(L + 1)].
130                                     activations,
131                                     delta_y = delta3)
132         else:
133             self.layers[str(L)].backprop(lambda x: self.diff_actFun(x, type=self
134                                     .actFun_type),
135                                     prev_activations = self.layers[str(L + 1)].
136                                     activations,
137                                     post_errors = self.layers[str(L - 1)].error,
138                                     post_weights = self.layers[str(L - 1)].
139                                     weights)
140
141     return None
142
143 def fit_model(self, X, y, epsilon=0.01, num_passes=20000, print_loss=True):
144     '''
145     fit_model uses backpropagation to train the network
146     :param X: input data
147     :param y: given labels
148     :param num_passes: the number of times that the algorithm runs through the
149     whole dataset
150     :param print_loss: print the loss or not
151     :return:
152     '''
153     # Gradient descent.
154     for i in range(0, num_passes):
155         # Forward propagation
156         self.feedforward(X, lambda x: self.actFun(x, type=self.actFun_type))
157         # Backpropagation
158         self.backprop(X, y)
159
160         # Add regularization terms (b1 and b2 don't have regularization terms)

```



```

154         for L in np.arange(1, self.num_layers):
155             self.layers[str(L)].dW += self.reg_lambda * self.layers[str(L)].
                weights
156
157             # Gradient descent parameter update
158             self.layers[str(L)].weights += -epsilon * self.layers[str(L)].dW
159             self.layers[str(L)].biases += -epsilon * self.layers[str(L)].db
160
161             # Optionally print the loss.
162             # This is expensive because it uses the whole dataset, so we don't want
                to do it too often.
163             if print_loss and i % 1000 == 0:
164                 print("Loss after iteration %i: %f" % (i, self.calculate_loss(X, y))
                    )
165
166     def visualize_decision_boundary(self, X, y):
167         '''
168         visualize_decision_boundary plots the decision boundary created by the
                trained network
169         :param X: input data
170         :param y: given labels
171         :return:
172         '''
173         plot_decision_boundary(lambda x: self.predict(x), X, y)
174
175     class Layer(object):
176         def __init__(self, layer_size, prec_layer_size, num_examples, seed, depth):
177             self.num_examples = num_examples
178             self.depth = depth
179             self.layer_size = layer_size
180             self.prec_layer_size = prec_layer_size
181             self.linear_output = np.zeros((num_examples, self.layer_size))
182             self.activations = np.zeros(self.linear_output.shape)
183             self.error = np.zeros((self.num_examples, self.layer_size))
184
185             # initialize the weights and biases in the network
186             np.random.seed(seed)
187             self.weights = np.random.randn(self.prec_layer_size, self.layer_size) / np.
                sqrt(self.prec_layer_size)
188             self.biases = np.zeros((1, self.layer_size))
189
190             # differentials for weights and biases
191             self.dW = np.zeros(self.weights.shape)
192             self.db = np.zeros(self.biases.shape)
193
194         def backprop(self, diff_actFun, prev_activations,
195                     post_errors = None, post_weights = None, delta_y = None):
196             '''
197             returns derivative of the loss function with respect to param
198             '''
199
200             # compute deltas
201             if self.depth == 1:
202                 self.error = delta_y
203             else:
204                 self.error = diff_actFun(self.linear_output) * post_errors.dot(
                    post_weights.T)
205
206             # compute weight or bias differential as average over examples

```

```

207         for example in range(self.num_examples):
208             vec1 = prev_activations[example, :]
209             vec1.shape = (1, len(vec1))
210             vec2 = self.error[example, :]
211             vec2.shape = (1, len(vec2))
212             self.dW += vec1.T.dot(vec2)
213
214         self.dW /= self.num_examples
215         self.db = self.error.mean(axis = 0)
216
217     def feedforward(self, inputs, actFun):
218         self.linear_output = inputs.dot(self.weights) + np.tile(self.biases, (inputs
219             .shape[0], 1))
220         self.activations = actFun(self.linear_output)
221         return None
222
223 def main():
224     # generate and visualize Make-Moons dataset
225     X, y = generate_data()
226     # plt.scatter(X[:, 0], X[:, 1], s=40, c=y, cmap=plt.cm.Spectral)
227     # plt.show()
228     model = NeuralNetwork(layer_sizes=[2, 3, 4, 2], num_examples= 200, actFun_type='
229         Tanh')
230     # model.fit_model(X,y_hot)
231     aa = datetime.datetime.now().replace(microsecond=0)
232     model.fit_model(X,y)
233     bb = datetime.datetime.now().replace(microsecond=0)
234     print('total time for training (hours:min:sec)', bb - aa)
235     model.visualize_decision_boundary(X,y)
236     # print(np.shape(self.W1), np.shape(self.W2))
237
238 if __name__ == "__main__":
239     main()

```

We trained a network with 2 hidden layers, first of sizes 3 and 4, and then, with sizes 30 and 40. The results are below, in figures 6 and 7.

```

Loss after iteration 0: 0.865495
Loss after iteration 1000: 0.312576
Loss after iteration 2000: 0.301356
Loss after iteration 3000: 0.289616
Loss after iteration 4000: 0.273894
Loss after iteration 5000: 0.252445
Loss after iteration 6000: 0.226776
Loss after iteration 7000: 0.201108
Loss after iteration 8000: 0.179541
Loss after iteration 9000: 0.163622
Loss after iteration 10000: 0.152680
Loss after iteration 11000: 0.145359
Loss after iteration 12000: 0.140475
Loss after iteration 13000: 0.137190
Loss after iteration 14000: 0.134956
Loss after iteration 15000: 0.133417
Loss after iteration 16000: 0.132344
Loss after iteration 17000: 0.131587
Loss after iteration 18000: 0.131046
Loss after iteration 19000: 0.130655
total time for training (hours:min:sec) 0:00:47

```




Figure 6: ‘Tanh’ activation function and 2 hidden layers with sizes 3 and 4

```

Loss after iteration 0: 1.215800
Loss after iteration 1000: 0.291953
Loss after iteration 2000: 0.262270
Loss after iteration 3000: 0.230261
Loss after iteration 4000: 0.198338
Loss after iteration 5000: 0.172955
Loss after iteration 6000: 0.155583
Loss after iteration 7000: 0.144192
Loss after iteration 8000: 0.136598
Loss after iteration 9000: 0.131339
Loss after iteration 10000: 0.127539
Loss after iteration 11000: 0.124682
Loss after iteration 12000: 0.122462
Loss after iteration 13000: 0.120688
Loss after iteration 14000: 0.119245
Loss after iteration 15000: 0.118056
Loss after iteration 16000: 0.117069
Loss after iteration 17000: 0.116248
Loss after iteration 18000: 0.115564
Loss after iteration 19000: 0.114997
total time for training (hours:min:sec) 0:01:04

```




Figure 7: ‘Tanh’ activation function and 2 hidden layers with sizes 30 and 40

2 Training a simple Deep Convolutional Network on MNIST

2.a Build and train a 4-layer DCN

2.a.1 Build network

Here are our personal implementations:

```

1 def weight_variable(shape):
2     '''
3     Initialize weights
4     :param shape: shape of weights, e.g. [w, h ,Cin, Cout] where
5     w: width of the filters
6     h: height of the filters
7     Cin: the number of the channels of the filters
8     Cout: the number of filters
9     :return: a tensor variable for weights with initial values
10    '''
11    initial = tf.truncated_normal(shape, stddev=0.1)
12    W = tf.Variable(initial, name = "W")
13
14    return W

1 def bias_variable(shape):
2     '''
3     Initialize biases
4     :param shape: shape of biases, e.g. [Cout] where

```

```

5     Cout: the number of filters
6     :return: a tensor variable for biases with initial values
7     '''
8     initial = tf.constant(0.1, shape=shape)
9     b = tf.Variable(initial, name = "b")
10    return b

1 def conv2d(x, W):
2     '''
3     Perform 2-D convolution
4     :param x: input tensor of size [N, W, H, Cin] where
5     N: the number of images
6     W: width of images
7     H: height of images
8     Cin: the number of channels of images
9     :param W: weight tensor [w, h, Cin, Cout]
10    w: width of the filters
11    h: height of the filters
12    Cin: the number of the channels of the filters = the number of channels of
        images
13    Cout: the number of filters
14    :return: a tensor of features extracted by the filters, a.k.a. the results after
        convolution
15    '''
16    h_conv = tf.nn.conv2d(x, W, strides=[1, 1, 1, 1], padding='SAME')
17    return h_conv

1 def max_pool_2x2(x):
2     '''
3     Perform non-overlapping 2-D maxpooling on 2x2 regions in the input data
4     :param x: input data
5     :return: the results of maxpooling (max-marginalized + downsampling)
6     '''
7     h_max = tf.nn.max_pool(x, ksize=[1, 2, 2, 1],
8                             strides=[1, 2, 2, 1], padding='SAME')
9     return h_max

```

2.a.2 Set up training

The rest of our `dcn_mnist.py` file looks like this:

```

1 def main():
2     # Specify training parameters
3     result_dir = './results/' # directory where the results from the training are
        saved
4     max_step = 5500 # the maximum iterations. After max_step iterations, the
        training will stop no matter what
5
6     start_time = time.time() # start timing
7
8     # placeholders for input data and input labels
9     x = tf.placeholder(tf.float32, shape=[None, 784], name = "x")
10    y_ = tf.placeholder(tf.float32, shape=[None, 10], name = "y")
11
12    # reshape the input image
13    x_image = tf.reshape(x, [-1, 28, 28, 1])
14
15    # first convolutional layer
16    W_conv1 = weight_variable([5, 5, 1, 32])
17    b_conv1 = bias_variable([32])

```

```
18 h_conv1 = tf.nn.relu(conv2d(x_image, W_conv1) + b_conv1)
19 h_pool1 = max_pool_2x2(h_conv1)
20
21 # second convolutional layer
22 W_conv2 = weight_variable([5, 5, 32, 64])
23 b_conv2 = bias_variable([64])
24
25 h_conv2 = tf.nn.relu(conv2d(h_pool1, W_conv2) + b_conv2)
26 h_pool2 = max_pool_2x2(h_conv2)
27
28 # densely connected layer
29 W_fc1 = weight_variable([7 * 7 * 64, 1024])
30 b_fc1 = bias_variable([1024])
31
32 h_pool2_flat = tf.reshape(h_pool2, [-1, 7*7*64])
33 h_fc1 = tf.nn.relu(tf.matmul(h_pool2_flat, W_fc1) + b_fc1)
34
35 # dropout
36 keep_prob = tf.placeholder(tf.float32)
37 h_fc1_drop = tf.nn.dropout(h_fc1, keep_prob)
38
39 # softmax
40 W_fc2 = weight_variable([1024, 10])
41 b_fc2 = bias_variable([10])
42
43 y_conv = tf.matmul(h_fc1_drop, W_fc2) + b_fc2
44
45 # setup training
46
47 cross_entropy = tf.reduce_mean(
48     tf.nn.softmax_cross_entropy_with_logits(labels=y_, logits=
49         y_conv))
50 train_step = tf.train.AdamOptimizer(1e-4).minimize(cross_entropy)
51 correct_prediction = tf.equal(tf.argmax(y_conv, 1), tf.argmax(y_, 1))
52 accuracy = tf.reduce_mean(tf.cast(correct_prediction, tf.float32))
53
54 # Add a scalar summary for the snapshot loss.
55 tf.summary.scalar(cross_entropy.op.name, cross_entropy)
56
57 # Build the summary operation based on the TF collection of Summaries.
58 summary_op = tf.summary.merge_all()
59
60 # Add the variable initializer Op.
61 init = tf.global_variables_initializer()
62
63 # Create a saver for writing training checkpoints.
64 saver = tf.train.Saver()
65
66 # Instantiate a SummaryWriter to output summaries and the Graph.
67 summary_writer = tf.summary.FileWriter(result_dir, sess.graph)
68
69 # Run the Op to initialize the variables.
70 sess.run(init)
71
72 # run the training
73 for i in range(max_step):
74     batch = mnist.train.next_batch(50) # make the data batch, which is used in
75         the training iteration.
76         # the batch size is 50
```

```

75     if i%100 == 0:
76         # output the training accuracy every 100 iterations
77         train_accuracy = accuracy.eval(feed_dict={x:batch[0], y_:batch[1],
78             keep_prob: 1.0})
79         print("step %d, training accuracy %g"%(i, train_accuracy))
80
81         # Update the events file which is used to monitor the training (in this
82         # case,
83         # only the training loss is monitored)
84         summary_str = sess.run(summary_op, feed_dict={x: batch[0], y_: batch[1],
85             keep_prob: 0.5}) # pb here
86         summary_writer.add_summary(summary_str, i)
87         summary_writer.flush()
88
89         # save the checkpoints every 1100 iterations
90         if i % 1100 == 0 or i == max_step:
91             checkpoint_file = os.path.join(result_dir, 'checkpoint')
92             saver.save(sess, checkpoint_file, global_step=i)
93
94         train_step.run(feed_dict={x: batch[0], y_: batch[1], keep_prob: 0.5}) # run
95         one train_step
96
97     # print test error
98     print("test accuracy %g"%accuracy.eval(feed_dict={
99         x: mnist.test.images, y_: mnist.test.labels, keep_prob: 1.0}))
100
101     stop_time = time.time()
102     print('The training takes %f second to finish'%(stop_time - start_time))
103
104 if __name__ == "__main__":
105     main()

```

2.a.3 Run training

After training, the final test accuracy of this network is 98.87%, as shown in the figure below:

```

jupyter DCN Last Checkpoint: 11 minutes ago (unsaved changes)
File Edit View Insert Cell Kernel Widgets Help Trusted Python (tensorflow) O
step 0, training accuracy 0.14
step 100, training accuracy 0.84
step 200, training accuracy 0.9
step 300, training accuracy 0.94
step 400, training accuracy 0.92
step 500, training accuracy 0.98
step 600, training accuracy 0.92
step 700, training accuracy 0.98
step 800, training accuracy 0.98
step 900, training accuracy 0.94
step 1000, training accuracy 0.94
step 1100, training accuracy 0.98
step 1200, training accuracy 0.98
step 1300, training accuracy 1
step 1400, training accuracy 0.94
step 1500, training accuracy 0.96
step 1600, training accuracy 0.96
step 1700, training accuracy 0.98
step 1800, training accuracy 0.98
step 1900, training accuracy 0.98
step 2000, training accuracy 0.94
step 2100, training accuracy 1
step 2200, training accuracy 0.96
step 2300, training accuracy 1
step 2400, training accuracy 0.96
step 2500, training accuracy 0.98
step 2600, training accuracy 0.98
step 2700, training accuracy 1
step 2800, training accuracy 1
step 2900, training accuracy 1
step 3000, training accuracy 0.98
step 3100, training accuracy 1
step 3200, training accuracy 0.98
step 3300, training accuracy 0.98
step 3400, training accuracy 0.94
step 3500, training accuracy 1
step 3600, training accuracy 0.96
step 3700, training accuracy 0.98
step 3800, training accuracy 1
step 3900, training accuracy 0.96
step 4000, training accuracy 1
step 4100, training accuracy 0.92
step 4200, training accuracy 0.98
step 4300, training accuracy 1
step 4400, training accuracy 0.98
step 4500, training accuracy 1
step 4600, training accuracy 1
step 4700, training accuracy 0.98
step 4800, training accuracy 1
step 4900, training accuracy 0.98
step 5000, training accuracy 0.98
step 5100, training accuracy 1
step 5200, training accuracy 1
step 5300, training accuracy 0.98
step 5400, training accuracy 1
test accuracy 0.9887
The training takes 441.854706 second to finish

```

Figure 8: Console output of first training of DCN on MNIST dataset

2.a.4 Visualize training

The visualization available via TensorBoard at this point are:

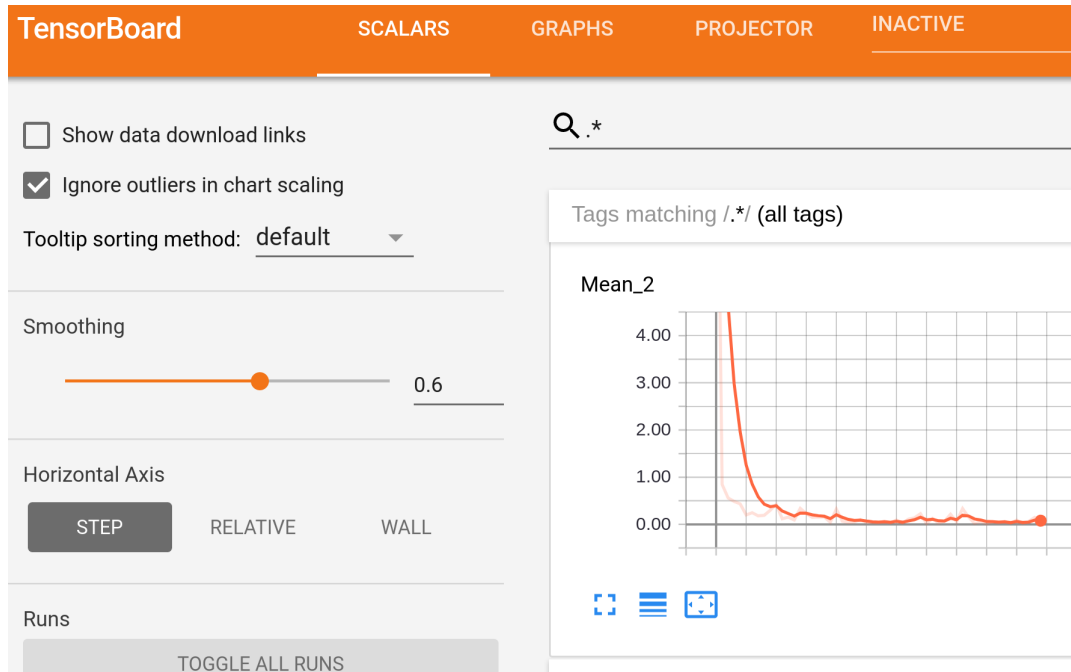


Figure 9: TensorBoard plot of loss as function of iterations, for first DCN training

Main Graph

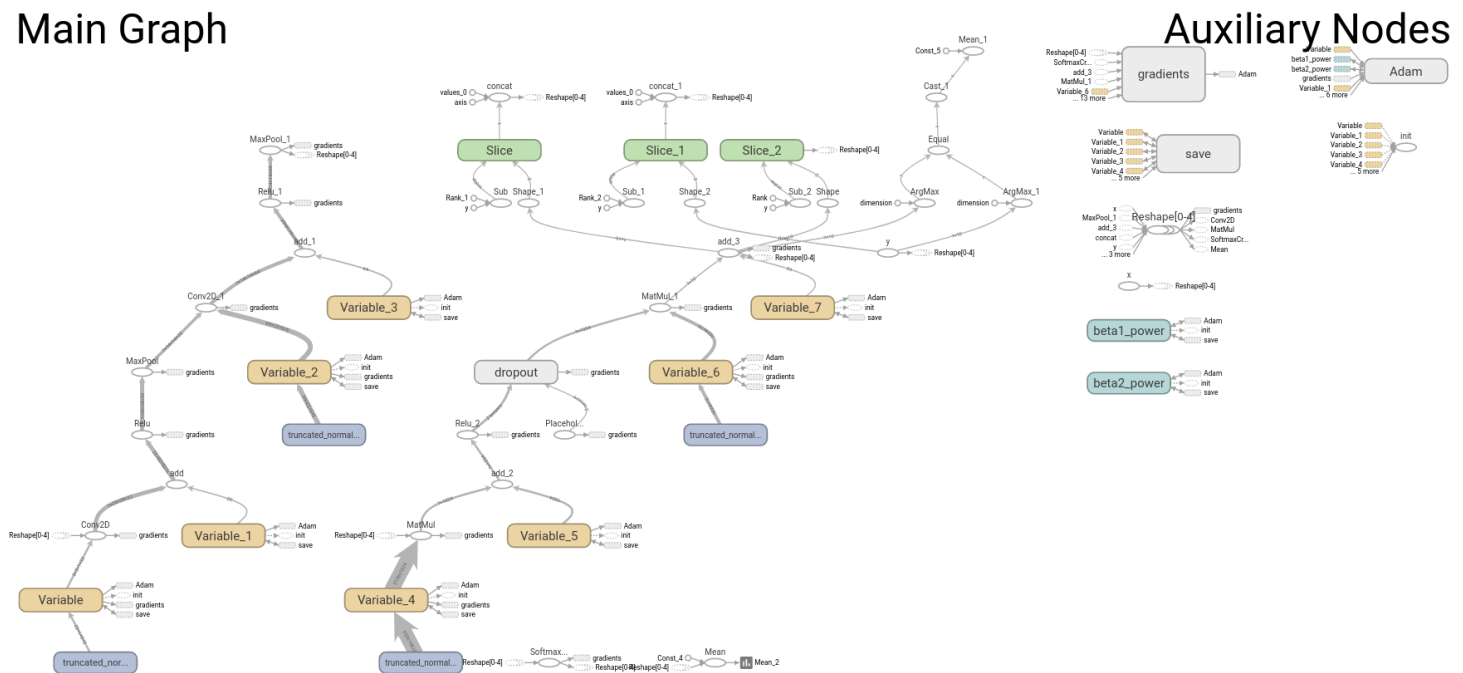


Figure 10: TensorBoard graph for first DCN training

2.b More on visualizing your training


After using namespaces and summaries in our code, the console output becomes: and TensorBoard graphs are:

2.c Time for more fun!

```
step 0, training accuracy 0.06
test accuracy 0.0981
validation accuracy 0.0954
step 100, training accuracy 0.8
step 200, training accuracy 0.8
step 300, training accuracy 0.82
step 400, training accuracy 0.92
step 500, training accuracy 0.98
step 600, training accuracy 0.94
step 700, training accuracy 0.98
step 800, training accuracy 0.92
step 900, training accuracy 0.96
step 1000, training accuracy 0.88
step 1100, training accuracy 0.94
test accuracy 0.9637
validation accuracy 0.9656
step 1200, training accuracy 0.96
step 1300, training accuracy 1
step 1400, training accuracy 0.94
step 1500, training accuracy 0.94
step 1600, training accuracy 0.96
step 1700, training accuracy 0.96
step 1800, training accuracy 1
step 1900, training accuracy 1
step 2000, training accuracy 0.98
step 2100, training accuracy 1
step 2200, training accuracy 1
test accuracy 0.9766
validation accuracy 0.9774
step 2300, training accuracy 0.98
step 2400, training accuracy 0.96
step 2500, training accuracy 1
step 2600, training accuracy 0.96
step 2700, training accuracy 1
step 2800, training accuracy 0.96
step 2900, training accuracy 0.98
step 3000, training accuracy 0.94
step 3100, training accuracy 0.96
step 3200, training accuracy 0.94
step 3300, training accuracy 0.98
test accuracy 0.9821
validation accuracy 0.9832
step 3400, training accuracy 0.98
step 3500, training accuracy 1
step 3600, training accuracy 1
step 3700, training accuracy 0.96
step 3800, training accuracy 1
step 3900, training accuracy 0.98
step 4000, training accuracy 1
step 4100, training accuracy 1
step 4200, training accuracy 1
step 4300, training accuracy 1
step 4400, training accuracy 0.98
test accuracy 0.9845
validation accuracy 0.9868
step 4500, training accuracy 0.98
step 4600, training accuracy 1
step 4700, training accuracy 1
step 4800, training accuracy 1
step 4900, training accuracy 0.98
step 5000, training accuracy 1
step 5100, training accuracy 0.98
step 5200, training accuracy 1
step 5300, training accuracy 1
step 5400, training accuracy 1
The training takes 467.104159 second to finish
```

Figure 11: Console output for second training of DCN

Main Graph



Auxiliary Nodes

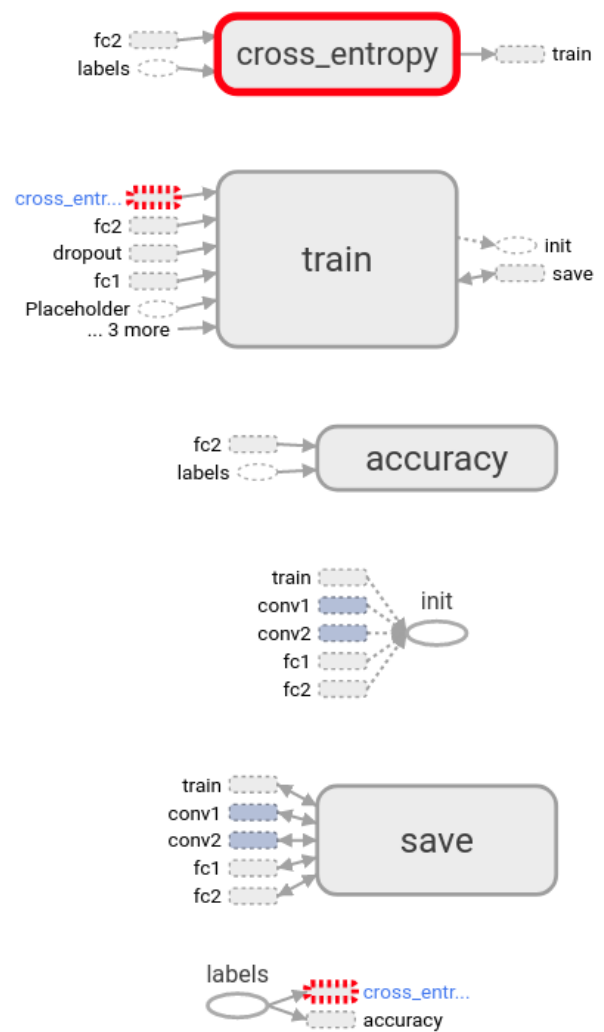



Figure 12: TensorBoard graph for second DCN training

accuracy_1



cross_entropy/Mean_1




Figure 13: TensorBoard plot of scalar summaries, for second DCN training

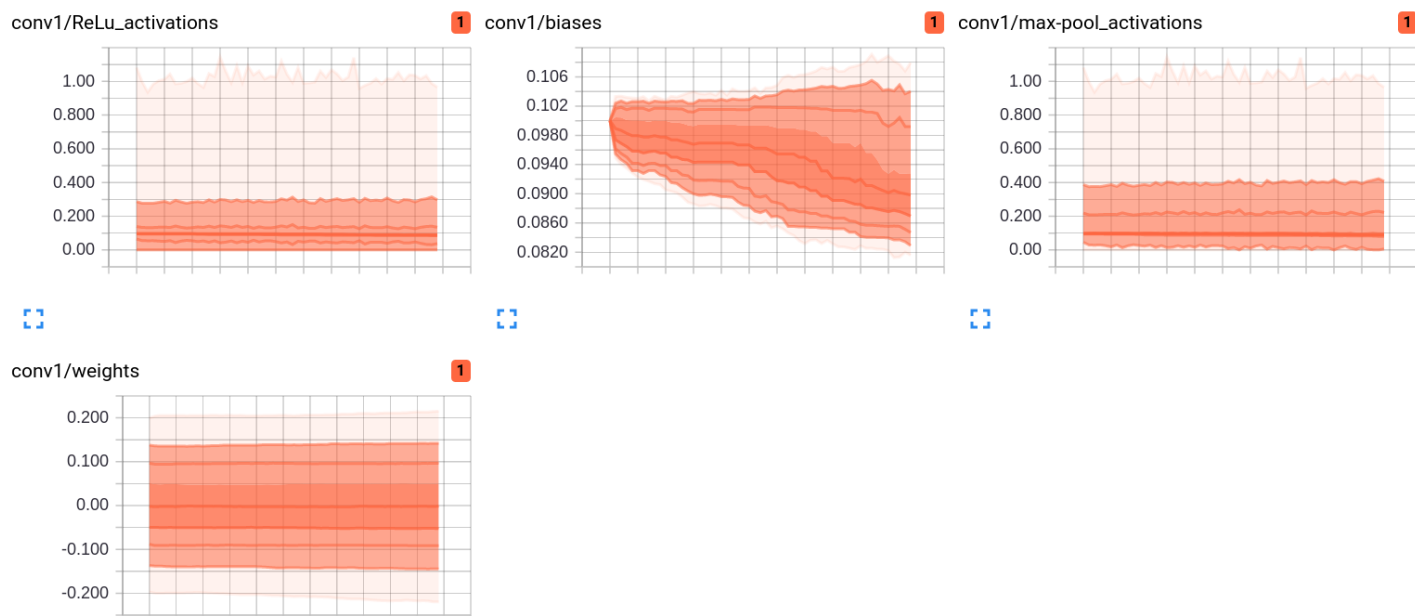


Figure 14: TensorBoard distributions in first convolutional layer, for second DCN training



Figure 15: TensorBoard distributions in second convolutional layer, for second DCN training




Figure 16: TensorBoard distributions in first fully connected layer, for second DCN training




Figure 17: TensorBoard distributions in output layer, for second DCN training

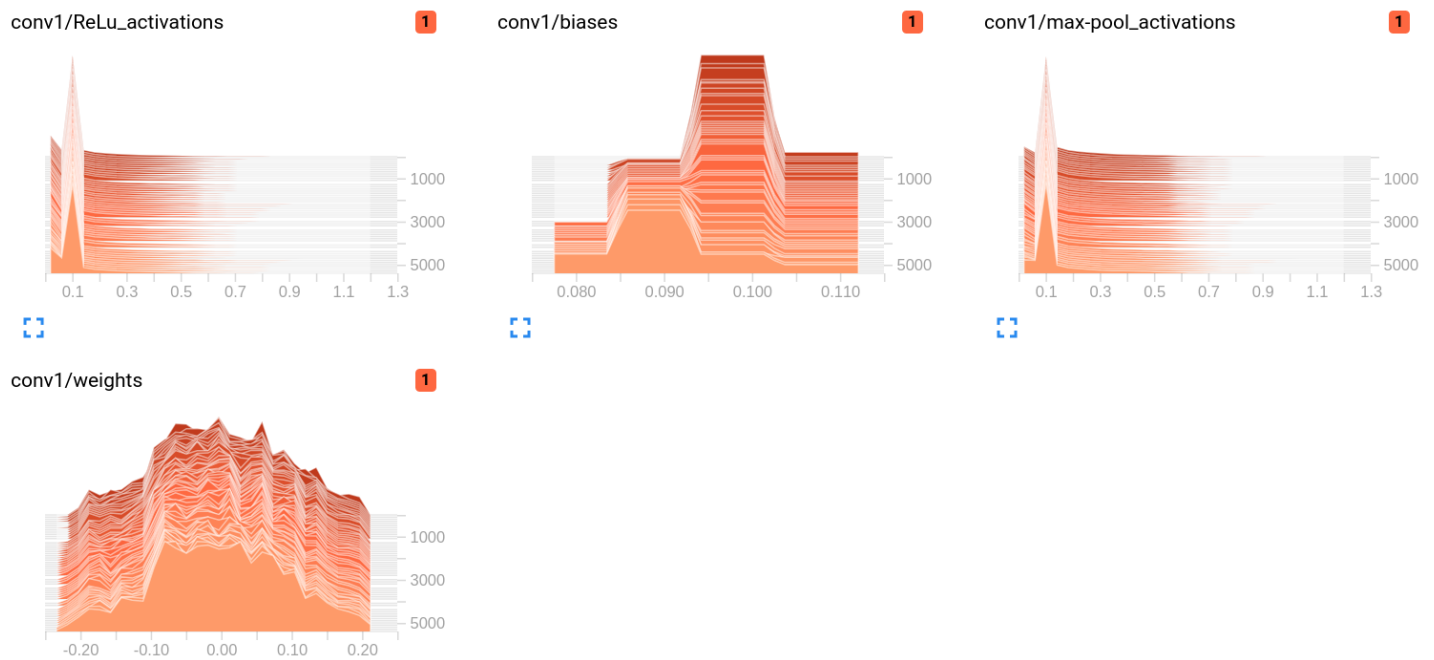


Figure 18: TensorBoard histograms in first convolutional layer, for second DCN training



Figure 19: TensorBoard histograms in second convolutional layer, for second DCN training




Figure 20: TensorBoard histograms in first fully connected layer, for second DCN training




Figure 21: TensorBoard histograms in output layer, for second DCN training